

Using Wasserstein GAN (with RL) for Natural Language Generation

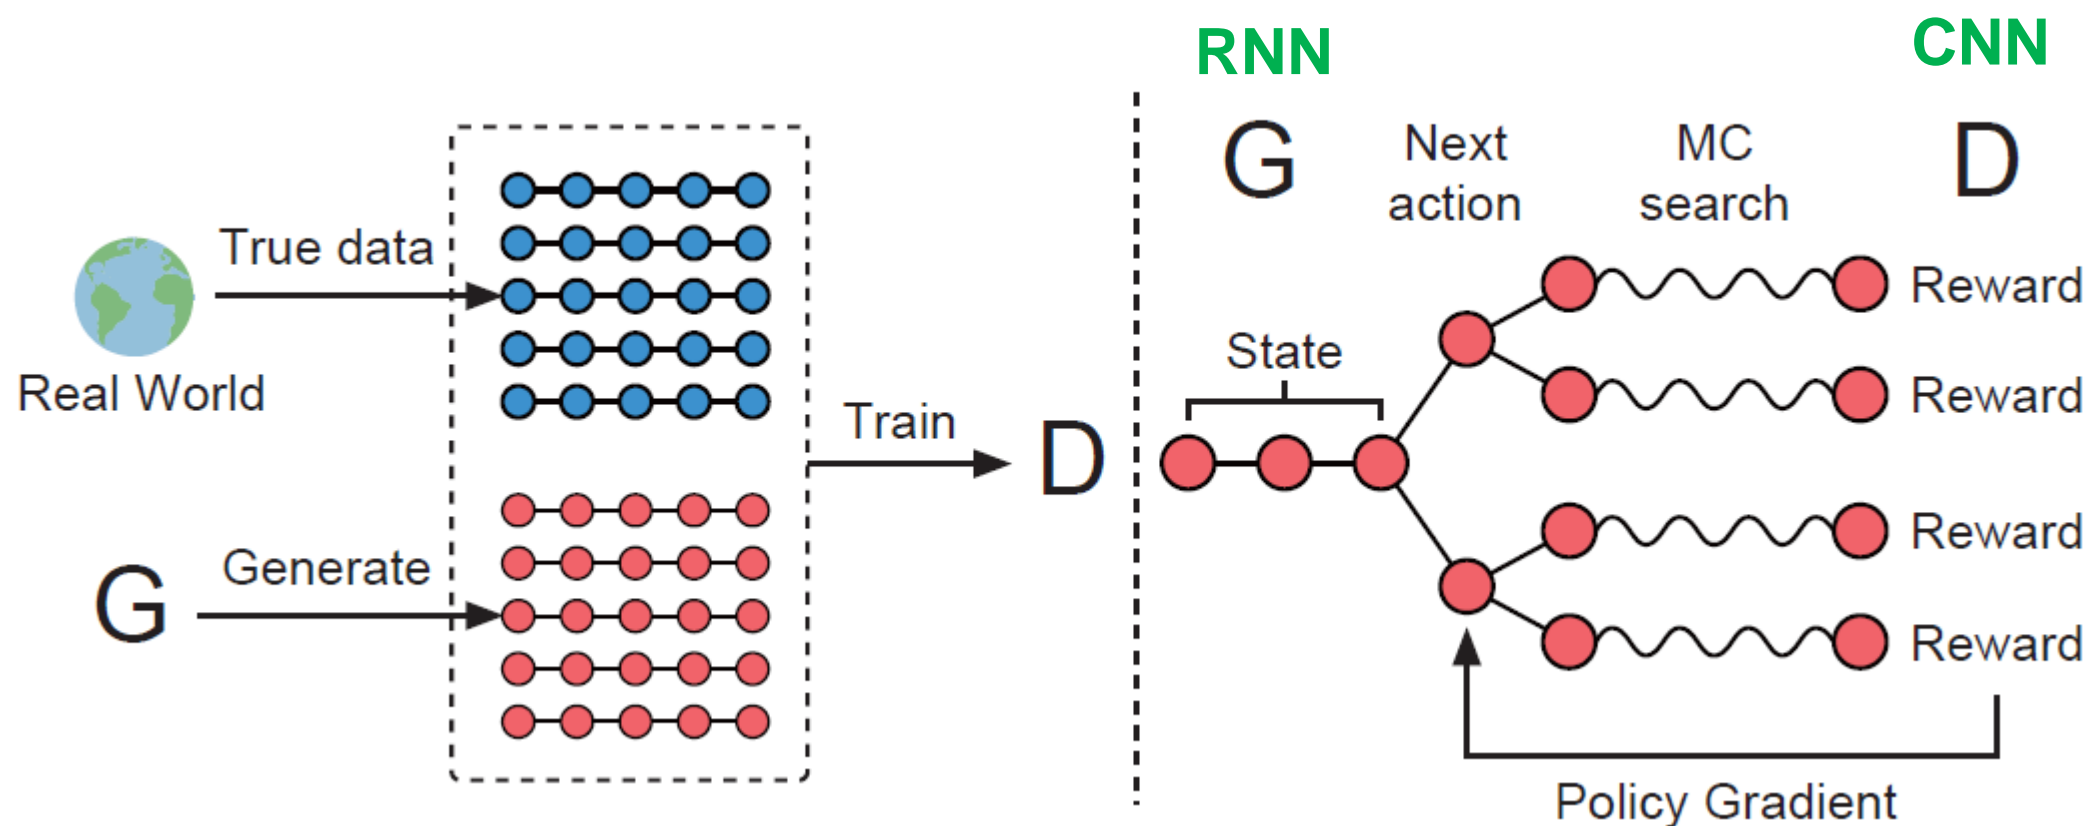
Guan-Yuan Chen



SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient

AAAI-17

Lantao Yu[†], Weinan Zhang^{†*}, Jun Wang[‡], Yong Yu[†]



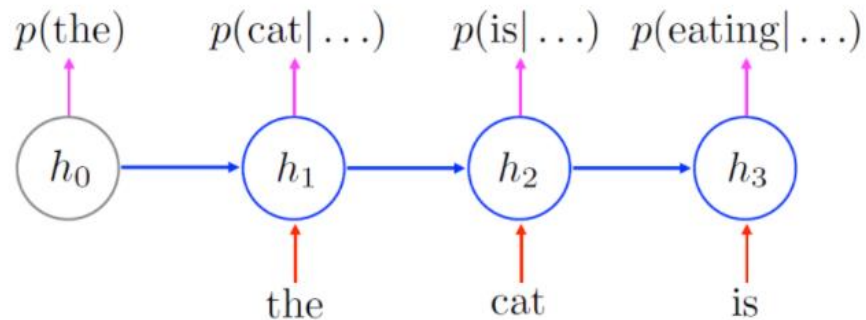
The problem of GAN on NLG

A major reason lies in that the **discrete outputs** from the generative model make it difficult to pass the gradient update from the discriminative model to the generative model.

The problem of GAN on RNN sampling

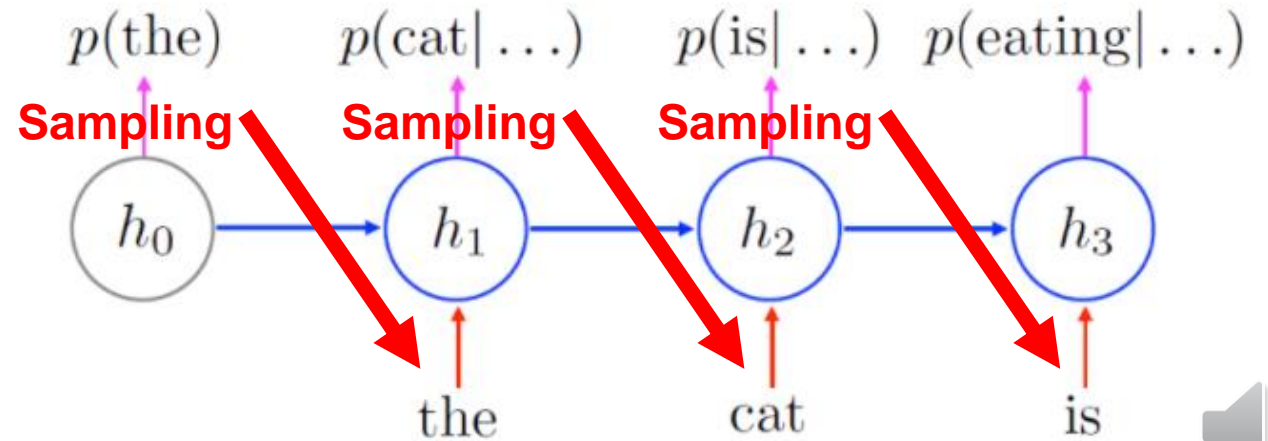
Training

$p(\text{the, cat, is, eating})$



Sampling


$p(\text{the, cat, is, eating})$



SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient

AAAI-17

Lantao Yu[†], Weinan Zhang^{†*}, Jun Wang[‡], Yong Yu[†]



Algorithm 1 Sequence Generative Adversarial Nets

Require: generator policy G_θ ; roll-out policy G_β ; discriminator D_ϕ ; a sequence dataset $\mathcal{S} = \{X_{1:T}\}$

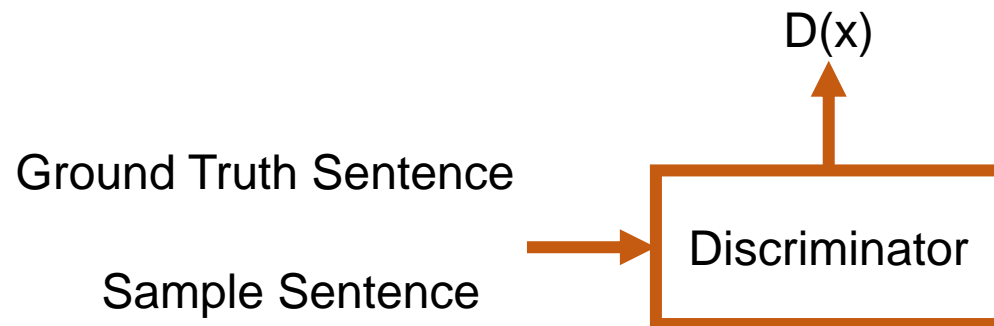
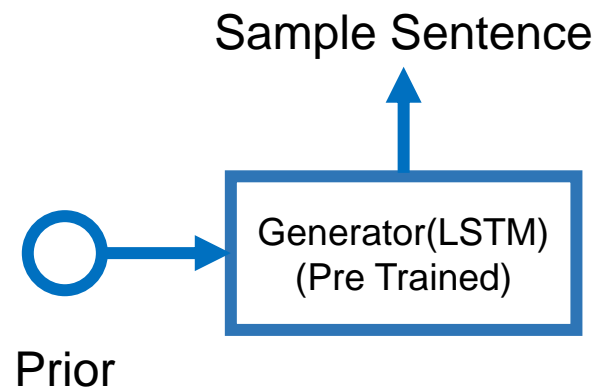
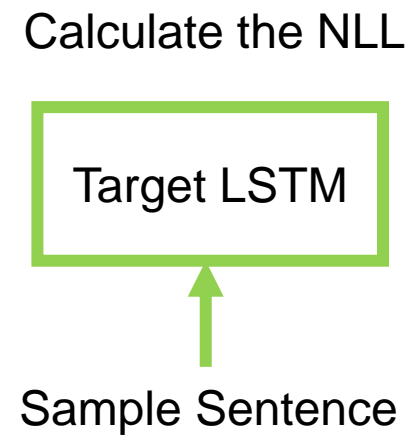

- 1: Initialize G_θ, D_ϕ with random weights θ, ϕ .
 - 2: Pre-train G_θ using MLE on \mathcal{S}
 - 3: $\beta \leftarrow \theta$
 - 4: Generate negative samples using G_θ for training D_ϕ
 - 5: Pre-train D_ϕ via minimizing the cross entropy
 - 6: **repeat**
 - 7: **for** g-steps **do**
 - 8: Generate a sequence $Y_{1:T} = (y_1, \dots, y_T) \sim G_\theta$
 - 9: **for** t in $1 : T$ **do**
 - 10: Compute $Q(a = y_t; s = Y_{1:t-1})$ by Eq. (4)
 - 11: **end for**
 - 12: Update generator parameters via policy gradient Eq. (8)
 - 13: **end for**
 - 14: **for** d-steps **do**
 - 15: Use current G_θ to generate negative examples and combine with given positive examples \mathcal{S}
 - 16: Train discriminator D_ϕ for k epochs by Eq. (5)
 - 17: **end for**
 - 18: $\beta \leftarrow \theta$
 - 19: **until** SeqGAN converges
-

$$Q_{D_\phi}^{G_\theta}(s = Y_{1:t-1}, a = y_t) = \begin{cases} \frac{1}{N} \sum_{n=1}^N D_\phi(Y_{1:T}^n), & Y_{1:T}^n \in \text{MC}^{G_\beta}(Y_{1:t}; N) & \text{for } t < T \\ D_\phi(Y_{1:t}) & & \text{for } t = T, \end{cases} \quad (4)$$

$$\nabla_\theta J(\theta) = \mathbb{E}_{Y_{1:t-1} \sim G_\theta} \left[\sum_{y_t \in \mathcal{Y}} \nabla_\theta G_\theta(y_t | Y_{1:t-1}) \cdot Q_{D_\phi}^{G_\theta}(Y_{1:t-1}, y_t) \right]. \quad (6)$$

$$\theta \leftarrow \theta + \alpha_h \nabla_\theta J(\theta), \quad (8)$$


$$\min_{\phi} -\mathbb{E}_{Y \sim p_{\text{data}}} [\log D_\phi(Y)] - \mathbb{E}_{Y \sim G_\theta} [\log(1 - D_\phi(Y))]. \quad (5)$$



Adversarial Training Step:

1. Update D
2. Copy the Generator be the Rollout LSTM
3. Update G by Policy Gradient
4. Calculate the NLL by Target LSTM


Algorithm	Random	MLE	SS	PG-BLEU	SeqGAN
NLL	10.310	9.038	8.985	8.946	8.736
p -value	$< 10^{-6}$	$< 10^{-6}$	$< 10^{-6}$	$< 10^{-6}$	




Problems of GAN:

Mode collapse problem,

The loss of G and D represent nothing (about the convergence)



Algorithm	Random	MLE	SS	PG-BLEU	SeqGAN
NLL	10.310	9.038	8.985	8.946	8.736
p -value	$< 10^{-6}$	$< 10^{-6}$	$< 10^{-6}$	$< 10^{-6}$	



$$\text{JSD}(P \parallel Q) = \frac{1}{2}D(P \parallel M) + \frac{1}{2}D(Q \parallel M)$$

where $M = \frac{1}{2}(P + Q)$



Earth Mover's Distance

Wasserstein GAN

Martin Arjovsky¹, Soumith Chintala², and Léon Bottou^{1,2}

2017

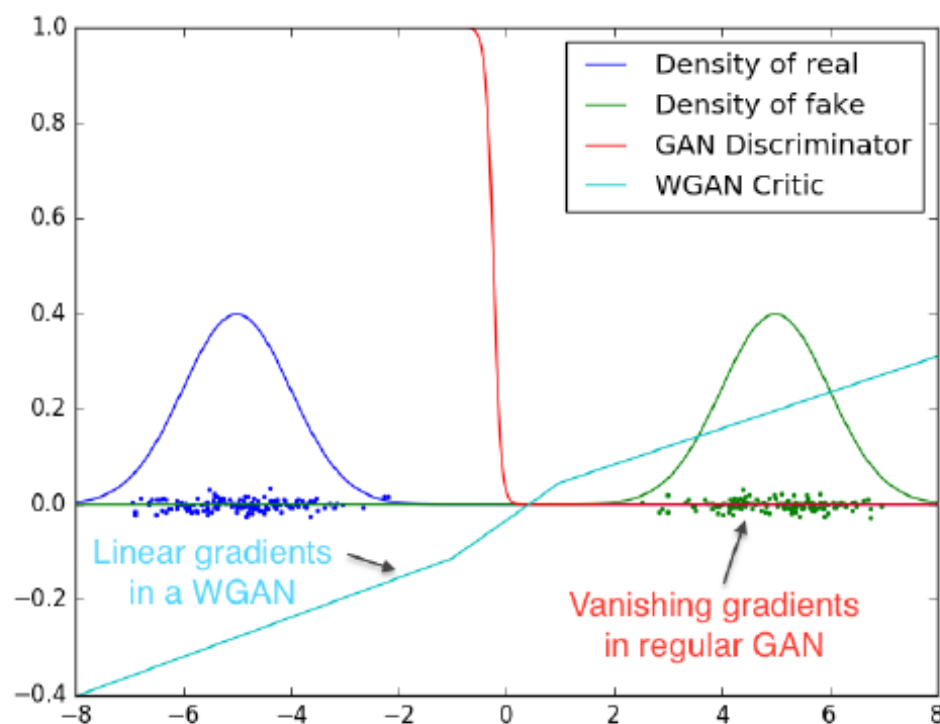
Wasserstein GAN

Generative Adversarial Nets

$$\max_{\|f\|_L \leq 1} \mathbb{E}_{x \sim \mathbb{P}_r} [f(x)] - \mathbb{E}_{x \sim \mathbb{P}_\theta} [f(x)]$$

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

$$\nabla_\theta W(\mathbb{P}_r, \mathbb{P}_\theta) = -\mathbb{E}_{z \sim p(z)} [\nabla_\theta f(g_\theta(z))]$$



Wasserstein GAN

Algorithm 1 WGAN, our proposed algorithm. All experiments in the paper used the default values $\alpha = 0.00005$, $c = 0.01$, $m = 64$, $n_{\text{critic}} = 5$.

Require: : α , the learning rate. c , the clipping parameter. m , the batch size. n_{critic} , the number of iterations of the critic per generator iteration.



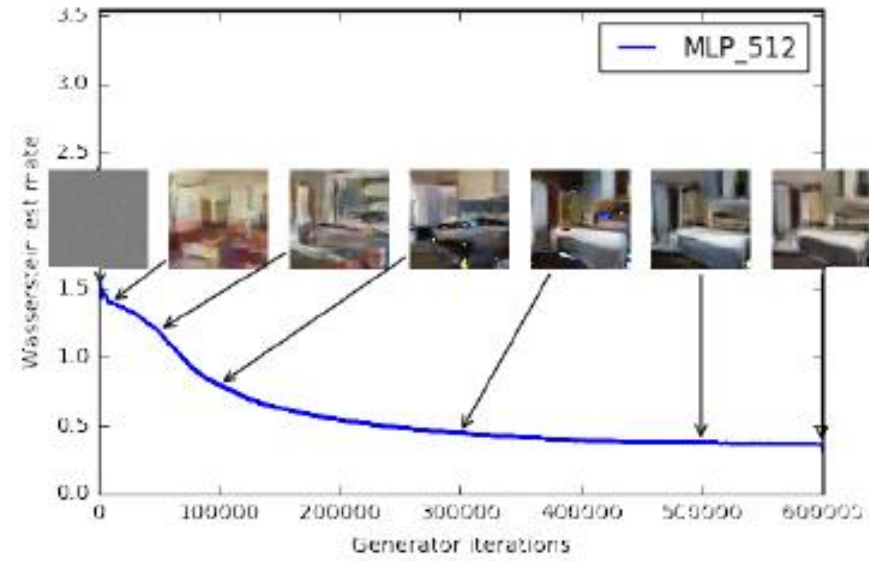
Require: : w_0 , initial critic parameters. θ_0 , initial generator's parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{\text{critic}}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w \left[ \frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSProp}(\theta, g_\theta)$ 
12: end while
```

Improved Training of Wasserstein GANs

Ishaan Gulrajani¹, Faruk Ahmed¹, Martin Arjovsky²

2017



Using Wasserstein GAN (with RL) for Natural Language Generation

方法:

以character-based的方式建構language generator(GRU)。

以RNN(GRU)的方式建構discriminator。

分別以Original GAN and Wasserstein GAN with Policy Gradient 之方式，訓練discriminator與generator。

Dataset: 7200 abstracts of Arxiv papers about neural networks (from Stanford CS 20SI)

$$J^{(D)} = -E_{x \sim p_{data}} [D(x)] + E_{z \sim p_z} [D(G(z))] \quad (1)$$

$$\theta_d = \theta_d - \eta \text{RMSP}rop(\theta_d, g_{\theta_d}), \quad \theta_d \leftarrow \text{clip}(\theta_d, c, -c) \quad (2)$$

其中， θ_d 為 discriminator 的參數。另定義 a prior on input noise variables $p_z(z)$ (from a random normal distribution)，作為 generator function 初始的 input value。細部的 procedure 與 algorithm 之定義如下：

Algorithm Wasserstein GAN with RL

Require: generator G_{θ_g} ; discriminator D_{θ_d} ; a sequence dataset $S = \{X_{1:T}\}$; roll-out policy G_ϕ

1: Initialize $G_{\theta_g}, D_{\theta_d}$ with random weights θ_g, θ_d ;

2: $\phi \leftarrow \theta_g$

3: Generate negative samples using G_{θ_g} for training D_{θ_d}

4: **repeat**

5: Pre-train G_{θ_g} using MLE on S

6: **for** g-steps **do**

7: Generate a sequence $Y_{1:T} = (y_1, \dots, y_T) \sim G_{\theta_g}$

8: **for** t in 1 : T **do**

9: Compute $Q(a = y_t; s = Y_{1:t-1})$ by Eq. (4) in [2]

10: **end for**

11: Update generator parameters via policy gradient Eq. (8) in [2]

12: **end for**

13: **for** d-steps **do**

14: Use current G_{θ_g} to generate negative examples and combine with given positive example S

15: Train discriminator D_{θ_d} for 5 epochs by Eq. (1) and (2)

16: **end for**

17: $\phi \leftarrow \theta_g$

18: **until** D_{θ_d} converges

Result

RNN Stack: 3, Hidden Size of G: 200, Hidden Size of D: 200	
SeqGAN	WGAN with PG, clipping value: 5
<p>=====</p> <p>the activation of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties</p> <p>=====</p>	<p>=====</p> <p>in neural networks in machine learning. One of techniques have yielded</p> <p>recognition, type. We study the conventional layer-wise pre-training to</p> <p>the individual models and deep neural networks as part of a method for</p> <p>stochastic gradient and analyzing algorithms. This parallel learning</p> <p>architectures are generic and deep neural network architectures. The</p> <p>gradients. In this paper, we propose a novel approach for nonlinear unit,</p> <p>to the recently proposed deep RNNs benefit</p> <p>=====</p>

Result

RNN Stack: 3, Hidden Size of G: 200, Hidden Size of D: 200	
SeqGAN	WGAN with PG, clipping value: 5
<p>=====</p> <p>the activation of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties of the network are a difficulties of the network are a</p> <p>difficulties</p> <p>=====</p>	<p>=====</p> <p>a fixed-point of computational complexity. Thus, we propose a new</p> <p>synchronization protocol that can be model distributed operation recent</p> <p>to train leventigation and dimensionality reduction in a data of noness of</p> <p>our methods on the same datasets in the relevant features of acoustic</p> <p>framework for example, representations of the approaches with a</p> <p>computationally expensive to train and deep learning architecture can</p> <p>be trained using standard points that by</p> <p>=====</p>

Result

RNN Stack: 3, Hidden Size of G: 200, Hidden Size of D: 200

SeqGAN

WGAN with PG, clipping value: 5

Discriminator Loss



Mean NLL




Name	Smoothed	Value	Step	Time	Relative
evod.	1.477	1.477	9.950k	Tue Jun 6, 17:29:41	2h 44m 39s

Name	Smoothed	Value	Step	Time	Relative
evod.	1.508	1.508	9.950k	Tue Jun 6, 20:23:46	2h 45m 51s

Result: Discriminator loss

RNN Stack 3, Hidden size of G 200, D 200

Adam

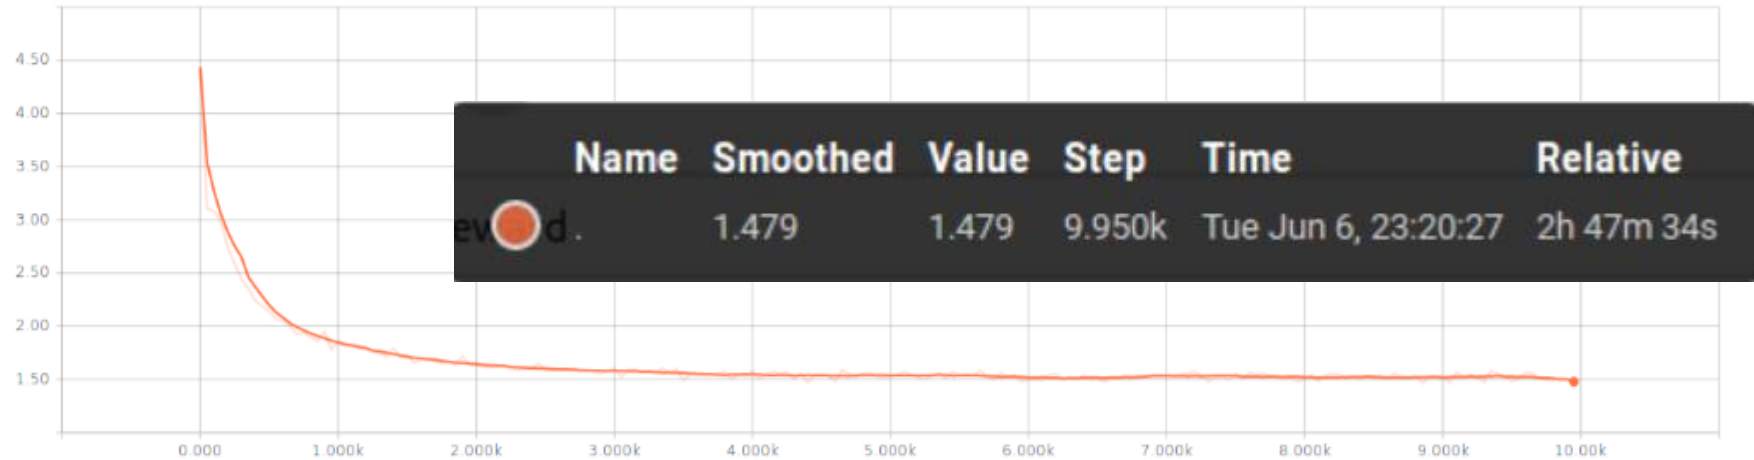


RMSprop

Result: Mean NLL

RNN Stack 3, Hidden size of G 200, D 200

Adam



RMSprop



References :

[1] Martin Arjovsky, Soumith Chintala, and Leon Bottou
Wasserstein GAN
2017

[2] Lantao Yuy, Weinan Zhangy, Jun Wangz, Yong Yu
SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient
2017

[3] Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, Aaron Courville
Improved Training of Wasserstein GANs
2017

[4] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza
Generative Adversarial Nets
2014

[5] <https://vincentherrmann.github.io/blog/wasserstein/>
Wasserstein GAN and the Kantorovich-Rubinstein Duality

[6] <https://github.com/LantaoYu/SeqGAN>
Implementation of Sequence Generative Adversarial Nets with Policy Gradient

[7] <http://blog.aylien.com/introduction-generative-adversarial-networks-code-tensorflow/>
An introduction to Generative Adversarial Networks (with code in TensorFlow)

[8] <http://web.stanford.edu/class/cs20si/>
CS 20SI: Tensorflow for Deep Learning Research